

Frank-Wolfe Algorithm Variants for LASSO Regression: A Comparative Study on High-Dimensional Datasets

Andrea Manocchio

andrea.manocchio@studenti.unipd.it

Giovanni Passaro

giovanni.passaro@studenti.unipd.it

September 13, 2025

Abstract

The Frank-Wolfe algorithm and its variants have gained renewed interest for solving constrained optimization problems in machine learning, particularly for sparse regression tasks. This paper presents a comprehensive comparison of three Frank-Wolfe variants—Vanilla Frank-Wolfe, Away-Steps Frank-Wolfe (AFW), and Pairwise Frank-Wolfe (PFW)—applied to LASSO regression problems. We evaluate these algorithms on diverse datasets including gene expression data (Leukemia), cancer genomics (TCGA), and medical prediction tasks (Heart Disease, Breast Cancer). Our empirical results demonstrate that Away-Steps Frank-Wolfe consistently outperforms other variants, achieving superior convergence rates and more effective sparse feature selection across all tested datasets. The study provides practical insights into algorithm selection for high-dimensional sparse optimization problems.

1 Introduction

The LASSO (Least Absolute Shrinkage and Selection Operator) regression problem is fundamental in modern machine learning and statistics, particularly for high-dimensional data analysis where feature selection is crucial. The problem seeks to minimize a quadratic objective subject to an ℓ_1 -norm constraint:

$$\min_{x \in \mathbb{R}^n} \|Ax - b\|_2^2 \quad \text{subject to} \quad \|x\|_1 \leq \tau \quad (1)$$

where $A \in \mathbb{R}^{m \times n}$ is the data matrix, $b \in \mathbb{R}^m$ is the target vector, and $\tau > 0$ is the regularization parameter controlling sparsity. Key characteristics are:

- **Automatic feature selection:** The L1 penalty forces some coefficients exactly to zero, eliminating irrelevant features
- **Overfitting control:** Particularly effective in high-dimensional settings where $p \gg n$
- **Interpretability:** Produces parsimonious and easily interpretable models

Mathematical Properties:

- **Convex optimization:** Guarantees global optimum, unlike non-convex alternatives
- **Bias-variance trade-off:** Introduces bias to significantly reduce variance

- **Shrinkage effect:** Coefficients are shrunk toward zero, with smaller effects eliminated entirely

Practical Considerations:

- **Parameter tuning:** Regularization parameter τ typically selected via cross-validation
- **Feature scaling:** Requires standardization since penalty is not scale-invariant
- **Multicollinearity:** Tends to arbitrarily select one feature from groups of highly correlated features

Applications in High-Dimensional Data:

- **Genomics:** Gene expression analysis and biomarker discovery
- **Medical imaging:** Feature selection in radiomics and pathology
- **Finance:** Risk factor identification and portfolio optimization
- **Text mining:** Document classification and sentiment analysis

Traditional optimization methods for this problem often require expensive projections onto the ℓ_1 -ball at each iteration. The Frank-Wolfe algorithm framework offers an attractive alternative by replacing projections with linear minimization oracle (LMO) calls, which are computationally efficient for structured constraint sets like the ℓ_1 -ball.

While the classical Frank-Wolfe algorithm is known to suffer from slow sublinear convergence due to zigzagging behavior near the constraint boundary, several variants have been developed to address these limitations. This paper investigates three prominent variants: the standard Frank-Wolfe algorithm, Away-Steps Frank-Wolfe (AFW), and Pairwise Frank-Wolfe (PFW).

Our contribution is a systematic empirical comparison of these algorithms on diverse high-dimensional datasets representative of modern machine learning applications, including genomics and medical prediction tasks.

2 Frank-Wolfe Algorithm Variants

2.1 Standard Frank-Wolfe Algorithm

The classical Frank-Wolfe algorithm iteratively moves toward vertices of the feasible set that minimize a linear approximation of the objective function. At iteration t , the algorithm:

1. Solves the linear subproblem: $s_t = \arg \min_{s \in \mathcal{D}} \langle \nabla f(x^{(t)}), s \rangle$
2. Computes the update direction: $d_t = s_t - x^{(t)}$
3. Determines step size γ_t (typically $\gamma_t = \frac{2}{t+2}$ or via line search)
4. Updates: $x^{(t+1)} = x^{(t)} + \gamma_t d_t$

For the ℓ_1 -ball constraint, the LMO has a simple closed-form solution: find the coordinate with maximum absolute gradient and set the corresponding vertex coordinate to $\pm\tau$ with appropriate sign.

2.2 Away-Steps Frank-Wolfe (AFW)

AFW addresses the zigzagging problem by introducing the ability to move away from vertices in the current active set. At each iteration, the algorithm chooses between a standard Frank-Wolfe step and an "away step" based on which offers greater progress:

1. Compute FW direction: $d_t^{FW} = s_t - x^{(t)}$
2. Compute away direction: $d_t^A = x^{(t)} - v_t$, where $v_t = \arg \max_{v \in S^{(t)}} \langle \nabla f(x^{(t)}), v \rangle$
3. Choose direction offering more progress
4. Perform line search with appropriate step size bounds
5. Update active set and vertex weights

The away step enables the algorithm to reduce weights on poorly-performing vertices, leading to better convergence properties.

2.3 Pairwise Frank-Wolfe (PFW)

PFW unifies the forward and away moves into a single pairwise direction that directly transfers mass between vertices:

1. Find FW vertex s_t and away vertex v_t as in AFW
2. Set pairwise direction: $d_t = s_t - v_t$
3. Set maximum step size: $\gamma_{\max} = \alpha_{v_t}$ (weight of away vertex)
4. Perform line search and update weights

This approach often yields the sparsest solutions by actively pruning vertices at each iteration.

3 Experimental Setup

3.1 Datasets

We evaluate the algorithms on four diverse high-dimensional datasets:

- **Leukemia:** Gene expression data with 7,129 features and 72 samples for cancer classification
- **TCGA Cancer:** Cancer genomics data with over 20,000 primary cancer and matched normal samples spanning 33 cancer types and engineered features for tumor classification
- **Heart Disease:** Medical prediction with clinical risk factors
- **Breast Cancer Wisconsin:** Cancer diagnosis with tumor characteristics

All features were standardized to zero mean and unit variance, and targets were centered. The regularization parameter τ was tuned for each dataset to ensure meaningful sparse solutions.

3.2 Implementation Details

All algorithms were implemented in Python with NumPy. The standard Frank-Wolfe used diminishing step sizes $\gamma_t = \frac{2}{t+2}$, while AFW and PFW employed exact line search. For AFW and PFW, we maintained explicit active sets and updated vertex weights according to the theoretical formulations.

Convergence was measured using both objective function values and Frank-Wolfe duality gaps. All experiments used a maximum of 300 iterations with tolerance 10^{-6} .

4 Results and Analysis

4.1 Leukemia Dataset: A Case Study in High-Dimensional Gene Expression

The Leukemia dataset, with 7,129 gene expression features and only 72 samples, represents an extreme case of high-dimensional sparse optimization.




Figure 1: Convergence comparison on Leukemia dataset. Left: Objective function values. Right: Frank-Wolfe duality gaps. AFW (orange) achieves exponential convergence reaching near machine precision (10^{-11}), while PFW (green) plateaus immediately at suboptimal values around 10^1 , and Vanilla FW (blue) shows characteristic sublinear behavior.

The results show AFW achieving exponential convergence, with the objective function decreasing from 10^3 to below 10^{-10} and the duality gap following a similar exponential decay pattern. This exceptional performance demonstrates AFW’s ability to navigate the high-dimensional gene expression space effectively.

In stark contrast, PFW exhibits immediate stagnation, with both objective values and duality gaps plateauing around 10^1 after the first few iterations. This behavior exemplifies the “gamma going to zero” problem where step sizes become negligibly small due to active set weight distribution issues in high-dimensional settings. In high-dimensional spaces with many features, the algorithm tends to activate many vertices (sparse solutions require many coordinates).

The PFW stagnation occurs because:

$$\gamma_{\max} = \alpha_{v_t} \rightarrow 0 \quad \Rightarrow \quad \gamma_t \approx O(1/n)$$

Combined with high-dimensional weight fragmentation $\alpha_i = O(1/\sqrt{n})$, the progress per iteration becomes:

$$\|\mathbf{x}^{(t+1)} - \mathbf{x}^{(t)}\| \leq O(1/\sqrt{n}) \rightarrow 0$$

resulting in practical algorithmic stagnation.

4.2 Cross-Algorithm Performance Analysis

4.2.1 AFW Dominance Across Datasets




Figure 2: AFW convergence across all datasets. The Leukemia dataset (blue) shows the most dramatic convergence reaching 10^{-12} , followed by Breast Cancer (purple) and Heart Disease (red) achieving convergence to 10^{-3} range. TCGA Cancer (orange) shows steady progress but remains in the 10^2 range due to higher dimensionality and complexity.

The convergence patterns reveal dataset-specific characteristics:

- **Leukemia:** Achieves the best convergence, reaching 10^{-12} due to strong linear separability in gene expression patterns
- **Breast Cancer and Heart Disease:** Show similar convergence profiles, reaching 10^{-3} range, indicating good problem conditioning
- **TCGA Cancer:** More challenging convergence remaining in 10^2 range, reflecting the complexity of multi-cancer genomics data

4.2.2 PFW Failure Pattern




Figure 3: PFW performance across datasets showing universal early stagnation. All datasets plateau within the first 50 iterations: Leukemia and Breast Cancer (purple lines) around $10^5 - 10^6$, TCGA Cancer (orange) around 10^2 , and Heart Disease (red) around 10^1 . This demonstrates the fundamental limitations of PFW in high-dimensional settings.

The uniform stagnation pattern across datasets confirms that PFW’s poor performance is not dataset-specific but rather a fundamental algorithmic limitation in high-dimensional sparse optimization.

4.2.3 Vanilla Frank-Wolfe Baseline Performance




Figure 4: Vanilla Frank-Wolfe convergence showing characteristic sublinear behavior. Leukemia (blue) achieves the best performance reaching 10^{-1} , while other datasets show slower convergence with more pronounced zigzagging effects. The duality gaps reveal the theoretical $O(1/k)$ convergence rate across all datasets.

While Vanilla FW shows predictable sublinear convergence, it significantly outperforms PFW while remaining far inferior to AFW, establishing a clear performance hierarchy.

4.3 Feature Selection and Sparsity Analysis

4.3.1 TCGA Cancer Dataset: Gene Selection Patterns

The TCGA cancer genomics dataset provides insights into how different algorithms approach feature selection in cancer biomarker discovery.

The sparsity analysis reveals:

- **AFW:** Selected 289 genes (1.4% sparsity) while achieving the lowest objective value, suggesting identification of a comprehensive cancer gene signature
- **Vanilla FW:** Selected 140 genes (0.7% sparsity) with intermediate performance
- **PFW:** Extremely sparse with only 17 genes (0.08% sparsity) but poor objective values, indicating premature convergence rather than effective feature selection

The gene overlap matrix shows that AFW’s superior performance stems from exploring a fundamentally different region of the gene expression space, identifying cancer biomarkers not found by other methods.

4.3.2 Heart Disease: Medical Risk Factor Identification

The medical risk factor analysis shows AFW identifying clinically meaningful features while maintaining superior convergence. The balanced identification of both risk factors (positive coefficients) and protective factors (negative coefficients) suggests AFW’s ability to capture the full complexity of cardiovascular disease prediction.

4.4 Feature Selection Quality Assessment

To better understand the sparsity patterns, we analyze the coefficient value distributions and feature overlap patterns across algorithms.

The analysis reveals that AFW not only converges better but also identifies features with stronger predictive signals, as evidenced by the higher magnitude coefficients in the selected feature sets.

5 Discussion

5.1 Algorithmic Performance Insights

Our evaluation reveals several critical insights:

5.1.1 AFW’s Success Factors

The consistent superiority of AFW can be attributed to:

1. **Effective zigzag mitigation:** Away steps prevent entrapment in suboptimal constraint boundary regions
2. **Robust active set management:** Maintains effective step sizes throughout optimization
3. **High-dimensional scalability:** Performance improves rather than degrades with increasing dimensionality

5.1.2 PFW’s High-Dimensional Limitations

The dramatic failure of PFW in our experiments highlights a critical gap between theoretical promise and practical performance:

1. **Step size degradation:** Pairwise step sizes become vanishingly small in high-dimensional active sets
2. **Premature convergence:** Algorithm stagnates far from optimal solutions
3. **Sparsity-performance trade-off:** Extreme sparsity comes at the cost of solution quality

5.2 Domain-Specific Applications

5.2.1 Genomics and Cancer Research

For cancer genomics applications, AFW’s ability to identify comprehensive gene signatures while maintaining convergence makes it particularly valuable for biomarker discovery. The identification of 289 cancer-related genes in the TCGA dataset, compared to PFW’s 17, could represent the difference between a clinically useful gene panel and an incomplete biomarker set.

5.2.2 Medical Insights

In medical applications where both accuracy and interpretability are crucial, AFW’s balanced approach to feature selection (identifying both risk and protective factors) provides clinically actionable insights while maintaining statistical rigor, while both FW and AFW are similarly good at identifying relevant features.




Figure 5: Both methods provide similar feature selection, while PFW underperforms on this task

6 Conclusions

This comprehensive empirical study found evidence for the superiority of Away-Steps Frank-Wolfe in high-dimensional LASSO regression applications. Key findings include:

1. **Convergence superiority:** AFW achieves exponential convergence rates across all datasets, with the Leukemia dataset demonstrating convergence to machine precision
2. **Feature selection effectiveness:** AFW identifies more comprehensive and higher-quality feature sets compared to alternatives
3. **High-dimensional robustness:** Performance scales favorably with increasing dimensionality

The study also reveals the significant limitations of Pairwise Frank-Wolfe in high-dimensional settings, challenging its theoretical appeal with practical performance concerns. The "gamma going to zero" phenomenon represents a fundamental barrier to PFW's application in modern machine learning contexts.

Away-Steps Frank-Wolfe emerges as a good compromise in both solution quality and interpretability. The algorithm's ability to navigate complex high-dimensional landscapes while maintaining meaningful sparsity makes it particularly valuable for scientific applications where feature selection has domain-specific interpretations.

Future research should investigate adaptive regularization strategies and hybrid approaches that could potentially combine AFW's convergence properties with enhanced sparsity control mechanisms.

References

- [1] M. Jaggi. Revisiting frank-wolfe: Projection-free sparse convex optimization. In *International conference on machine learning*, pages 427–435. PMLR, 2013.
- [2] S. Lacoste-Julien and M. Jaggi. On the global linear convergence of frank-wolfe optimization variants. *Advances in neural information processing systems*, 28, 2015.
- [3] I. M. Bomze, F. Rinaldi, and D. Zeffiro. Frank–wolfe and friends: a journey into projection-free first-order optimization methods. *4OR*, 19:313–345, 2021.
- [4] R. Tibshirani. Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society: Series B*, 58(1):267–288, 1996.
- [5] R. Tibshirani. <https://archive.ics.uci.edu/dataset/45/heart+disease>.
- [6] R. Tibshirani. <https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>.